



# Make hiring smarter.

An AI system that ranks candidates the way a great recruiter would — reasoning about who actually fits, not who has the most keywords.

**Redrob** • Intelligent Candidate Discovery & Ranking

Senior AI Engineer JD • 100,000-candidate pool • Nimisha Agarwal

# Keyword filters can't see fit

The dataset is built with four traps. A keyword/embedding matcher walks straight into them.

1

## Keyword stuffers

Non-engineering profiles padded with AI skill tags. Look perfect on paper; can't do the job.

2

## Plain-language Tier-5s

Built a real recommendation system but never wrote "RAG" or "Pinecone." Easy to miss.

3

## Behavioral twins

Identical skills, opposite reachability — one responds to recruiters, one vanished months ago.

4

## Honeypots (~80)

Subtly impossible profiles. Ranking them in your top 10 means you aren't reading profiles.

# Rank on what the JD means, not what it says

A great recruiter reads between the lines. So does this system.

## WHAT IT SAYS

“5–9 years of experience”

A list of 20+ skills

“embeddings, vector DBs, NDCG”

Location: Pune / Noida

## WHAT IT MEANS

A range, not a cutoff — senior judgment


Shipped a ranking/search/recsys system

at a product company, to real users

Reachable & open — or it's not a hire

# One interpretable score

Every term lives in  $[0,1]$ , so the pieces read straight into each candidate's reasoning line.



Career evidence is weighted above the skills list (defeats stuffers, surfaces plain-language fits). The multipliers are the recruiter's judgment: the JD's explicit rejections, whether the person is actually available, and whether the profile is internally consistent.

**Pure standard-library Python — no model downloads, no API calls, fully auditable.**

# Read what they did, not what they tagged

The semantic signal is computed over career prose — summary + job descriptions — never the skills tags.

## How

TF-IDF cosine between each candidate's narrative and a JD concept-query (retrieval, ranking, recommendation, embeddings, vector search, NDCG, production, NLP/IR).

## Why prose, not tags

A “Marketing Manager” can list every AI keyword. They can't fake “owned search relevance and shipped a re-ranker to production.” Prose is where real fit — and plain-language Tier-5s — live.

**0.45**

weight on the semantic term

**1,141**


term corpus vocabulary

**prose**

summary + every job description

# A structured rubric carries the rest

55% of the base — domain evidence leads, role-fit guards against keyword stuffers.



## Skill-trust, not skill-count

Each core skill is weighted by  
**proficiency ×  $\sqrt{\text{duration}}$  × endorsements**

An “expert” skill with 0 months used and 0 endorsements counts as padding (×0.1) — so keyword stuffing earns almost nothing here.

# The JD's “do NOT want,” encoded

Each rejection is a multiplicative penalty — strong elsewhere can still partly recover.



Pure research, no production

×0.30



Entire career at IT-services / consulting

×0.28



CV / speech / robotics, little NLP-IR

×0.45



AI = recent LangChain/OpenAI wrapping only ×0.55



Senior who no longer writes code

×0.62




Title-chaser — many short stints

×0.78

# A perfect profile that's unreachable isn't a hire

Activity recency, response rate, open-to-work and notice period modulate every score.



## The JD is explicit

*"A perfect-on-paper candidate who hasn't logged in for 6 months and has a 5% response rate is, for hiring purposes, not actually available."*

A Meta Lead AI Engineer and a Microsoft Senior NLP Engineer both fall below the top-100 cutoff — on availability alone, not skill.



# Reading for consistency, not keywords

Honeypots are caught by internal contradictions — no special-casing, just careful reading.

1

Tenure exceeds experience — role months sum to more years than the stated career

2

Dates contradict duration — duration\_months disagrees with start/end by >18 months

3

Ghost “expert” skills — 4+ advanced/expert skills with 0 months and 0 endorsements

4

A single role longer than the whole career — impossible by construction

0 / 100

**honeypots in our top 100.**

Flagged profiles are multiplied by 0.02 and sink. (>10% in the top 100 is an automatic disqualification.)

# The shortlist a recruiter can trust

Top of the list: Zomato, Genpact AI, Razorpay, Paytm, Salesforce, Sarvam AI, CRED, Ola.



**0**  
non-engineering titles in the  
top 100

**88/100**  
inside the 5-9 year band

**47 + 51**  
in Tier-1 cities / willing to  
relocate

**100**  
unique, specific reasoning  
lines

# Production-grade by construction

Built for the constraint the JD actually cares about: scale, not a benchmark.

**0**

third-party deps in the ranking step

**CPU**

no GPU, no network, no API calls

**~2 min**

to rank all 100,000 candidates

**< 1 GB**

RAM — two streaming passes

**1**

command reproduces submission.csv

**stdlib**

fully auditable, trivially reproducible



# Understand the role. Read the career. Down-weight the unreachable.

That's how a great recruiter works — and it's exactly what this ranker does.

Delivered: [clean GitHub repo](#) • [top-100 submission.csv](#) • [runnable Colab sandbox](#)